

FRANK ALBANESE

Serial CTO & Co-Founder • AI/ML • Distributed Systems • Enterprise
Platforms

New York, NY [linkedin.com/in/frank-albanese-a8b17885](https://www.linkedin.com/in/frank-albanese-a8b17885) github.com/faalbane

March 23, 2026

Hiring Manager
Anthropic
New York, NY

Re: Solutions Architect, Applied AI (Industries)

Dear Hiring Manager,

I am writing to express my strong interest in the Solutions Architect, Applied AI (Industries) role at Anthropic. As a serial CTO, Y Combinator founder, and hands-on AI/ML builder with over a decade of experience delivering production AI systems, guiding enterprise customers through complex technical adoptions, and architecting scalable platforms, I am deeply aligned with Anthropic's mission to build reliable, interpretable, and steerable AI—and I am excited by the opportunity to help large enterprises successfully integrate and deploy Claude into their technology stacks.

My experience at FlutterFlow (Y Combinator W21) maps directly to the pre-sales and technical advisory work at the core of this role. As Enterprise Customer Success Engineer, I served as the senior technical partner for enterprise clients across industries, guiding them from initial technical discovery through successful deployment. I architected custom solutions integrating AI/ML capabilities, cloud-native services (Firebase, GCP, AWS), and third-party APIs for complex enterprise requirements. I drove product strategy feedback loops directly between enterprise customers and engineering leadership, influencing platform roadmap and feature prioritization—the exact cross-functional coordination between Sales, Product, and Engineering teams that this role demands. I developed and delivered technical content ranging from deep-dive architecture sessions for engineering teams to business value presentations for executive stakeholders, and I led enablement initiatives that established best practices for security, performance, and scalable application architecture.

As a two-time CTO and Y Combinator founder, I bring the deep technical expertise and customer-facing skills needed to architect innovative LLM solutions that

address complex business challenges. At CopyCatch.AI, I currently define and execute the technology vision for an AI-powered content authentication platform, architecting scalable AI/ML pipelines that combine LLM reasoning with domain-specific evaluation—the kind of applied AI architecture design that helps customers understand Claude's capabilities and build effective evaluation frameworks. At CoVerified (Y Combinator W21), I owned end-to-end AI product delivery in a highly regulated environment, designing distributed systems with compliance-first engineering and driving applied AI/ML from research to production. My four-plus years at Henry Ford Health System building ML-powered predictive tools and data-driven applications under HIPAA compliance further demonstrates my ability to deliver AI systems that organizations can trust in sensitive, enterprise contexts.

What excites me most about this role is the opportunity to sit at the intersection of Anthropic's frontier research and real-world enterprise deployment. I have spent my career translating cutting-edge AI capabilities into production-ready systems that deliver measurable outcomes—and I am passionate about helping customers navigate the technical architecture decisions, evaluation frameworks, and integration patterns that determine whether AI adoption succeeds at scale. I thrive in the kind of ambiguity inherent in pre-sales technical advisory work, where every customer presents a unique set of challenges, and the ability to rapidly understand requirements, design scalable architectures, and build trust as a technical partner is what separates good outcomes from great ones.

Anthropic's commitment to building AI that is safe and beneficial resonates deeply with my own values. The opportunity to help enterprises adopt Claude responsibly, to contribute integration insights back to Product and Engineering, and to advance the goal of safe, beneficial AI through real-world deployment is exactly the kind of high-impact work I am looking for. As a New York City resident comfortable with Python, cloud architectures, and cross-organizational collaboration, I am ready to contribute immediately.

I would welcome the opportunity to discuss how my experience in enterprise technical advisory, AI platform architecture, and customer-facing technical leadership can contribute to Anthropic's Applied AI mission. Thank you for your time and consideration.

Sincerely,

Frank Albanese